Chapter 1

Files with FileSystem

with the participation of:
Max Leske

Filesystem is the new file system library for Pharo. It is available in the system since Pharo 1.4. Filesystem has been developed by Colin Putney. This chapter is a quick start that shows how to get started.

1.1   Getting started

The framework supports different kinds of filesystems that can be used interchangeably and that can transparently work with each other. The most obvious one is the filesystem on your hard disk. We are going to work with that one for now. FSFilesystem is the factory to access different filesystem.

Sending the message disk to FSFilesystem, returns a file system as on your physical harddrive. Another less used possibility is memory to create a file system at the system of the image.

```
| working |
working := FSFilesystem disk workingDirectory.
```

The message workingDirectory above returns a reference to the working directory. References are instances of the class FSReference. As we will see references are the central objects of the framework and provide the primary mechanisms for working with files and directories.

Filesystem defines four classes that are important for the end-user: FSFilesystem, FSReference, FSLocator, and FSDirectoryEntry. You should do not use platform specific classes such as FSUnixStore or FSWindowsStore, there are internal classes. All code snippets below work on FSReference instances.

1.2   Navigating the Filesystem

Now let's do some more interesting things. To list the immediate children of your working directory, execute the following expression:

```
| working |
working := FSFilesystem disk workingDirectory.
working children.
```

Notice that children returns the files and folders. To recursively access all the children of the current directory you should use the message allChildren:

```
working allChildren.
```

To find all jpeg files in working, execute:

```
working allChildren select: [ :each | each basename endsWith: 'jpeg' ]
```

Use the slash operator to obtain a reference to a specific file or directory within your working directory:

```
| working cache |
working := FSFilesystem disk workingDirectory.
cache := working / 'package-cache'.
```

Navigating back to the parent is easy using the parent message:

```
parent := cache parent.
parent = working            ⟶     true
```

You can check for various properties of the cache directory by executing the following expressions:

```
cache exists.             ⟶     true
cache isFile.             ⟶     false
cache isDirectory.        ⟶     true
cache basename.           ⟶     'package-cache'
```

The methods exists, isFile, isDirectory, and basename are defined on the FSReference class.

To get additional information about a filesystem entry, we should get an FSDirectoryEntry using the message entry. Here are some examples:

```
cache entry creation.       ⟶    2010-02-14T10:34:31+00:00
cache entry modification.   ⟶    2010-02-14T10:34:31+00:00
cache entry size.           ⟶    0 (directories have size 0)
```

The framework also supports locations, late-bound references that point to a file or directory. When asking to perform a concrete operation, a location behaves the same way as a reference. Currently the following locations are supported:

```
FSLocator desktop.
FSLocator home.
FSLocator imageDirectory.
FSLocator image.
FSLocator changes.
FSLocator vmBinary.
FSLocator vmDirectory.
```

If you save a location with your image and move the image to a different machine or operating system, a location will still resolve to the expected directory or file. Note that some of them are still in flux because depending on specific VM functionalities.

Opening Read- and Write-Streams

Stéf ►*this example seems conceptually bogus with the current state of FS*◄ To open a file-stream on a file ask the reference for a read- or write-stream:

```
| working stream |
working := FSFilesystem disk workingDirectory.
stream := (working / 'foo.txt') writeStream.
stream nextPutAll: 'Hello World'.
stream close.
stream := (working / 'foo.txt') readStream.
stream contents.          ⟶     'Hello World'
stream close.
```

Stéf ►*in Pharo 1.4 we do not get 'Hello World' but a byte array*◄

Please note that writeStream overrides any existing file and readStream throws an exception if the file does not exist. There are also short forms available that eliminate the need to close a stream manually:

```
working / 'foo.txt' writeStreamDo: [ :stream | stream nextPutAll: 'Hello World' ].
working / 'foo.txt' readStreamDo: [ :stream | stream contents ].
```

```
FSFilesystem disk
   openFileStream: 'authors.txt' writable: true
```

```
(FSFilesystem disk workingDirectory / 'authors2.txt')
   writeStreamDo: [ :stream | stream nextPutAll: 'lkhklkjh']
```

Have a look at the streams protocol of FSReference for other convenience methods.

Renaming, Copying and Deleting Files and Directories

You can also copy and rename files by executing:

```
working / 'foo.txt' copyTo: (working / 'bar.txt').
working / 'foo.txt' renameAs: 'zork.txt'
```

> **Stéf** ▶ *renameAs: does not work since it calls unimplemented methods* ◀

To create a directory use the message createDirectory:

```
backup := working / 'cache−backup'.
backup createDirectory.
```

To then copy the contents of the complete package-cache to that directory simply use copyAllTo::

```
cache copyAllTo: backup.
```

Note that the target directory will be automatically created if it was not there before.

To delete a single file, use the message delete:

```
(working / 'bar.txt') delete.
```

To delete a complete directory tree use deleteAll. Be careful with that one though.

1.3   Design

> **Stéf** ▶ *should check class comments and a uml diagram* ◀

Now we explain the key classes of Filesystem.

Path

Paths are the most fundamental element of the Filesystem API. They represent filesystem paths in a very abstract sense, and provide a high-level protocol for working with paths without having to manipulate strings. Here are some examples showing how to define absolute paths (/), relative paths (*), file extension (,), parent navigation (parent)

"absolute path"
FSPath / 'plonk' / 'feep'   ⟶   /plonk/feep

"relative path"
FSPath * 'plonk' / 'feep'   ⟶   plonk/feep

"relative path with extension"
FSPath * 'griffle' , 'txt'   ⟶   griffle.txt

"changing the extension"
FSPath * 'griffle.txt' , 'jpeg'   ⟶   griffle.jpeg

"parent directory"
(FSPath / 'plonk' / 'griffle') parent   ⟶   /plonk

"resolving a relative path"
(FSPath / 'plonk' / 'griffle') resolve: (FSPath * '..' / 'feep')
              ⟶   /plonk/feep

"resolving an absolute path"
(FSPath / 'plonk' / 'griffle') resolve: (FSPath / 'feep')
              ⟶   /feep

"resolving a string"
(FSPath * 'griffle') resolve: 'plonk'   ⟶   griffle/plonk

"comparing"
(FSPath / 'plonk') contains: (FSPath / 'griffle' / 'nurp')
              ⟶   false

Note that some of the path protocol (messages like /, parent and resolve:) are also available on references — references are a combination of path and filesystem.

Filesystem

A filesystem is an interface to access hierarchies of directories and files. "The filesystem," provided by the host operating system, is represented by FSDiskFilesystem and it's platform-specific subclasses. However, the user should not access them directly but using FSFilesystem as we show previously. Other kinds of Filesystems are also possible. The memory filesystem provides a RAM disk filesystem where all files are stored as ByteArrays in the image. The zip filesystem represents the contents of a zip file.

Each filesystem has its own working directory, which it uses to resolve any relative paths that are passed to it. Some examples:

```
fs := FSFilesystem memory.
fs workingDirectoryPath: (FSPath / 'plonk').
griffle := FSPath / 'plonk' / 'griffle'.
nurp := FSPath * 'nurp'.

fs resolve: nurp.              ⟶      /plonk/nurp

fs createDirectory: (FSPath / 'plonk').   ⟶     "/plonk created"
(fs writeStreamOn: griffle) close.   ⟶     "/plonk/griffle created"
fs isFile: griffle.           ⟶     true
fs isDirectory: griffle.          ⟶     false
fs copy: griffle to: nurp.        ⟶     "/plonk/griffle copied to /plonk/nurp"
fs exists: nurp.             ⟶     true
fs delete: griffle.           ⟶     "/plonk/griffle" deleted
fs isFile: griffle.           ⟶     false
fs isDirectory: griffle.          ⟶     false
```

Reference

Paths and filesystems are the lowest level of the Filesystem API. A
FSReference combines a path and a filesystem into a single object which pro-
vides a simpler protocol for working with files. It implements the same op-
erations as FSFilesystem, but without the need to track paths and filesystem
separately:

```
fs := FSFilesystem memory.
griffle := fs referenceTo: (FSPath / 'plonk' / 'griffle').
nurp := fs referenceTo: (FSPath * 'nurp').

griffle isFile.           ⟶      false
griffle isDirectory.          ⟶      false

griffle parent ensureDirectory.
griffle writeStreamDo: [ :stream | ].
griffle exists & griffle isFile      ⟶      true
griffle copyTo: nurp
nurp exists.             ⟶      true
griffle delete
```

References also implement the path protocol with methods like /, parent
and resolve:.

Locator

Locators could be considered late-bound references. They're left deliberately
fuzzy, and are only resolved to a concrete reference when some file operation

needs to be performed. Instead of a filesystem and path, locators are made
up of an origin and a path. An origin is an abstract filesystem location, such
as the user's home directory, the image file, or the VM executable. When it
receives a message like isFile, a locator will first resolve its origin, then resolve
its path against the origin.

Locators make it possible to specify things like "an item named 'package-
cache' in the same directory as the image file" and have that specification
remain valid even if the image is saved and moved to another directory, pos-
sibly on a different computer.

```
locator := FSLocator image / 'package−cache'.
locator printString.          ⟶    '{image}/package−cache'
locator resolve.              ⟶    /Users/colin/Projects/Mason/package−cache
locator isFile.           ⟶    false
locator isDirectory.          ⟶    true
```

The following origins are currently supported:

- imageDirectory - the directory in which the image resides in

- image - the image file

- changes - the changes file

- vmBinary - the executable for the running virtual machine

- vmDirectory - the directory containing the VM application (may not be
  the parent of vmBinary)

- home - the user's home directory

- desktop - the directory that hold the contents of the user's desktop

- documents - the directory where the user's documents are stored (e.g.
  '/Users/colin/Documents')

Applications may also define their own origins, but the system will not
be able to resolve them automatically. Instead, the user will be asked to man-
ually choose a directory. This choice is then cached so that future resolution
requests will not require user interaction.

Enumeration

References and Locators also provide simple methods for dealing with
whole directory trees:

allChildren. This will answer an array of references to all the files and directories in the directory tree rooted at the receiver. If the receiver is a file, the array will contain a single reference, equal to the receiver.

allEntries. This method is similar to allChildren, but it answers an array of FSDirectoryEntries, rather than references.

copyAllTo: aReference. This will perform a deep copy of the receiver, to a location specified by the argument. If the receiver is a file, the file will be copied; if a directory, the directory and its contents will be copied recursively. The argument must be a reference that doesn't exist; it will be created by the copy.

deleteAll. This will perform a recursive delete of the receiver. If the receiver is a file, this has the same effect as delete.

Visitors

The above methods are sufficient for many common tasks, but application developers may find that they need to perform more sophisticated operations on directory trees.

The visitor protocol is very simple. A visitor needs to implement visitFile: and visitDirectory:. The actual traversal of the filesystem is handled by a guide. A guide works with a visitor, crawling the filesystem and notifying the visitor of the files and directories it discovers. There are three Guide classes, FSPreorderGuide, FSPostorderGuide and FSBreadthFirstGuide , which traverse the filesystem in different orders. To arrange for a guide to traverse the filesystem with a particular visitor is simple. Here's an example:

```
FSBreadthFirstGuide show: aReference to: aVisitor
```

The enumeration methods described above are implemented with visitors; see FSCopyVisitor, FSDeleteVisitor, and FSCollectVisitor for examples.